

An aerial photograph of a dense urban area, likely New York City, showing a grid of streets, tall buildings, and a mix of modern and older architecture. A prominent yellow arrow points diagonally across the image from the top right towards the bottom left. Overlaid on the bottom left is a large, semi-transparent teal graphic element that tapers to a point.

here

Deep Learning with Keras

Shubhabrata Roy, HERE Research Eindhoven
Pycon Ireland | November 06, 2016

Introduction to Deep Learning and Keras

here

1. What exactly is deep learning ?

**2. Why is it generally better than other methods on certain types of data
(on image, speech and some other types)?**

here

1. What exactly is deep learning ?

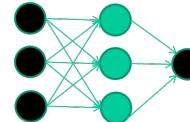
**2. Why is it generally better than other methods on certain types of data
(on image, speech and some other types)?**

**1. ‘Deep Learning’ implies using a neural network
with several layers of nodes between input and output**

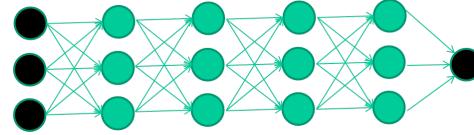
**2. The series of layers between input & output perform
feature identification and processing in a series of stages,
just as our brains do.**

Multilayer neural networks have been around for decades. What's actually new?

We have always had good algorithms for learning the weights in networks with 1 hidden layer



But these algorithms are not good at learning the weights for networks with more hidden layers



What's new is: algorithms for training many-layer networks

Introduction to Keras

- Developed by Francois Chollet from Google
- A minimalist Python library for deep learning
- Can run on top of Theano or TensorFlow
- Possible to quickly develop and test models
- Runs on both Python 2.7 and 3.5
- Can seamlessly execute on GPUs and CPUs

here

A bit more about Keras

Top libraries by Github forks

#1:	9004	tensorflow/tensorflow
#2:	6069	BVLC/caffe
#3:	1670	fchollet/keras
#4:	1367	Theano/Theano
#5:	1287	dmlc/mxnet
#6:	1252	torch/torch7
#7:	1076	Microsoft/CNTK
#8:	972	deeplearning4j/deeplearning4j
#9:	916	karpathy/convnetjs
#10:	585	Lasagne/Lasagne
#11:	376	NervanaSystems/neon
#12:	330	NVIDIA/DIGITS
#13:	312	tensorflow/skflow
#14:	281	pfnet/chainer
#15:	235	mila-udem(blocks
#16:	166	autumnai/leaf
#17:	148	tflearn/tflearn
#18:	99	IDSLA/brainstorm

Top libraries by Github contributors

#1:	227	Theano/Theano
#2:	222	tensorflow/tensorflow
#3:	193	BVLC/caffe
#4:	174	fchollet/keras
#5:	129	dmlc/mxnet
#6:	96	torch/torch7
#7:	72	deeplearning4j/deeplearning4j
#8:	64	Microsoft/CNTK
#9:	54	pfnet/chainer
#10:	46	Lasagne/Lasagne
#11:	45	mila-udem(blocks
#12:	38	NervanaSystems/neon
#13:	24	NVIDIA/DIGITS
#14:	23	tensorflow/skflow
#15:	14	karpathy/convnetjs
#16:	14	autumnai/leaf
#17:	13	IDSLA/brainstorm
#18:	11	tflearn/tflearn

Top libraries by Github issues opened

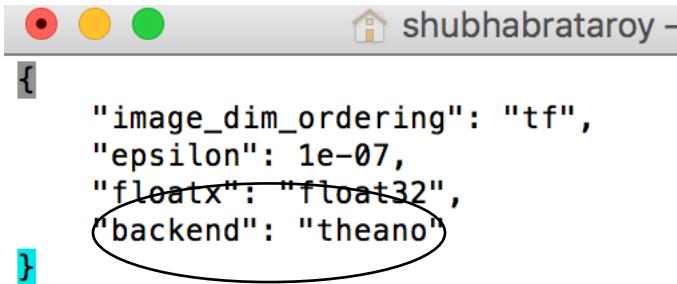
#1:	2587	BVLC/caffe
#2:	1946	fchollet/keras
#3:	1556	tensorflow/tensorflow
#4:	1537	Theano/Theano
#5:	1212	dmlc/mxnet
#6:	823	deeplearning4j/deeplearning4j
#7:	484	mila-udem(blocks
#8:	419	Microsoft/CNTK
#9:	412	pfnet/chainer
#10:	393	NVIDIA/DIGITS
#11:	357	Lasagne/Lasagne
#12:	304	torch/torch7
#13:	195	NervanaSystems/neon
#14:	105	tensorflow/skflow
#15:	81	IDSLA/brainstorm
#16:	56	tflearn/tflearn
#17:	43	autumnai/leaf
#18:	41	karpathy/convnetjs

Aggregate popularity ($30 \cdot \text{contrib} + 10 \cdot \text{issues} + 5 \cdot \text{forks}$) $\cdot 10^{-3}$

#1:	67.24	tensorflow/tensorflow
#2:	62.01	BVLC/caffe
#3:	33.03	fchollet/keras
#4:	29.02	Theano/Theano
#5:	22.43	dmlc/mxnet
#6:	15.25	deeplearning4j/deeplearning4j
#7:	12.18	torch/torch7
#8:	11.49	Microsoft/CNTK
#9:	7.88	Lasagne/Lasagne
#10:	7.37	mila-udem(blocks
#11:	7.15	pfnet/chainer
#12:	6.30	NVIDIA/DIGITS
#13:	5.41	karpathy/convnetjs
#14:	4.97	NervanaSystems/neon
#15:	3.30	tensorflow/skflow
#16:	1.70	IDSLA/brainstorm
#17:	1.68	autumnai/leaf
#18:	1.63	tflearn/tflearn

Theano and TensorFlow Backends

- A lightweight API
- Provides a consistent interface to efficient numerical libraries called backends
- Easiest way to choose the backend is editing `~/.keras/keras.json`



```
{  
    "image_dim_ordering": "tf",  
    "epsilon": 1e-07,  
    "floatx": "float32",  
    "backend": "theano"  
}
```

Building deep learning models with Keras

- **Define the model:** Create a Sequential model and add configured layers
- **Compile the model:** Specify loss function and optimizers and call the compile() function on the model
- **Fit the model:** Train the model on a sample of data by calling the fit() function on the model
- **Make predictions:** Use the model to generate predictions on new data by calling functions such as evaluate() or predict() on the model

Available layers in Keras

Regular Dense/ Multi Layer Perceptron

```
keras.layers.core.Dense(output_dim,  
                        init='glorot_uniform',  
                        activation='linear',  
                        weights=None,  
                        W_regularizer=None, b_regularizer=None, activity_regularizer=None,  
                        W_constraint=None, b_constraint=None,  
                        input_dim=None)
```

Recurrent layer/ LSTM/ GRU etc.

```
keras.layers.recurrent.GRU(output_dim,  
                           init='glorot_uniform', inner_init='orthogonal',  
                           activation='sigmoid', inner_activation='hard_sigmoid',  
                           return_sequences=False,  
                           go_backwards=False,  
                           stateful=False,  
                           input_dim=None, input_length=None)
```

More layers

1-D Convolutional Layers

```
keras.layers.convolutional.Convolution1D(nb_filter, filter_length,  
    init='uniform',  
    activation='linear',  
    weights=None,  
    border_mode='valid',  
    subsample_length=1,  
    W_regularizer=None, b_regularizer=None,  
    W_constraint=None, b_constraint=None,  
    input_dim=None, input_length=None)
```

2-D Convolutional Layers

```
keras.layers.convolutional.Convolution2D(nb_filter, nb_row, nb_col,  
    init='glorot_uniform',  
    activation='linear',  
    weights=None,  
    border_mode='valid',  
    subsample=(1, 1),  
    W_regularizer=None, b_regularizer=None,  
    W_constraint=None,  
    dim_ordering='th')
```

Other types of layer include:

- Dropout
- Noise
- Pooling
- Normalization
- Embedding

Activations

- Almost all major activations are available: Sigmoid, tanh, ReLu, softplus, hard sigmoid, linear
- Advanced activations implemented as a layer (after desired neural layer)
- Advanced activations: LeakyReLu, PReLU, ELU, Parametric Softplus, Thresholded linear and Thresholded Relu

here

Objectives and Optimizers

Objective Functions:

- Error loss: rmse, mse, mae, mape, msle
- Hinge loss: squared hinge, hinge
- Class loss: binary crossentropy, categorical crossentropy

Optimization:

- Provides SGD, Adagrad, Adadelta, Rmsprop and Adam
- All optimizers can be customized via parameters

A brief introduction to “HERE”

HERE in numbers

200
Countries mapped

7,000+

Employees in 56 countries focused on
delivering the world's best map and
location services

Millions of
changes made to the map
every day

Today, HERE provides services to nearly all global OEM brands as well as many market leading brands such Microsoft, Amazon, Facebook etc.

30+

Years of experience
transforming mapping
technology

4 of 5

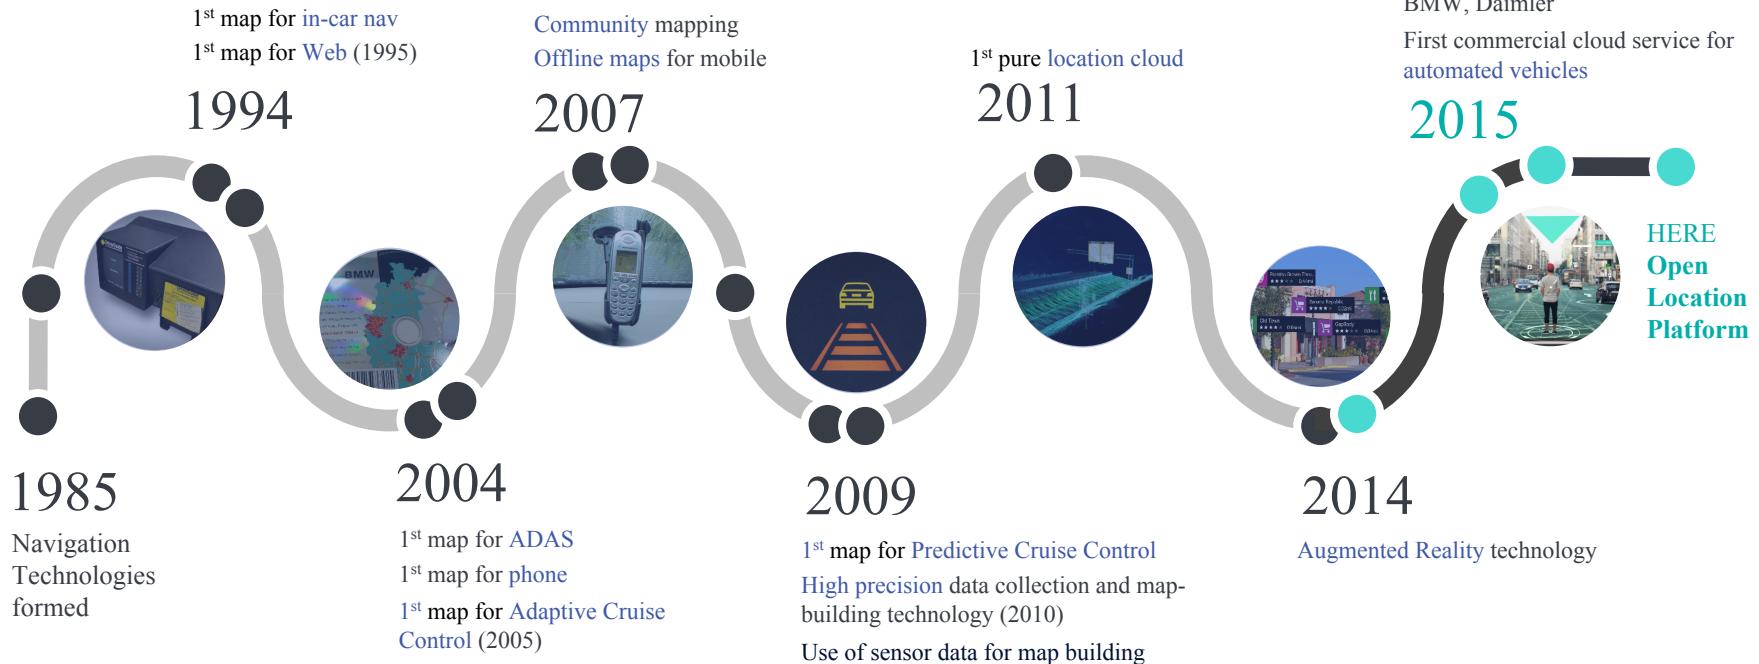
In-car navigation systems in Europe and
N America use HERE maps



We are making sense of the world through the lens of location.

We help people achieve better outcomes –
whether it's a driver moving to their destination safely,
a municipality managing its infrastructure smartly or
a business optimizing the utility of its assets.

A history of transforming potential into meaningful products



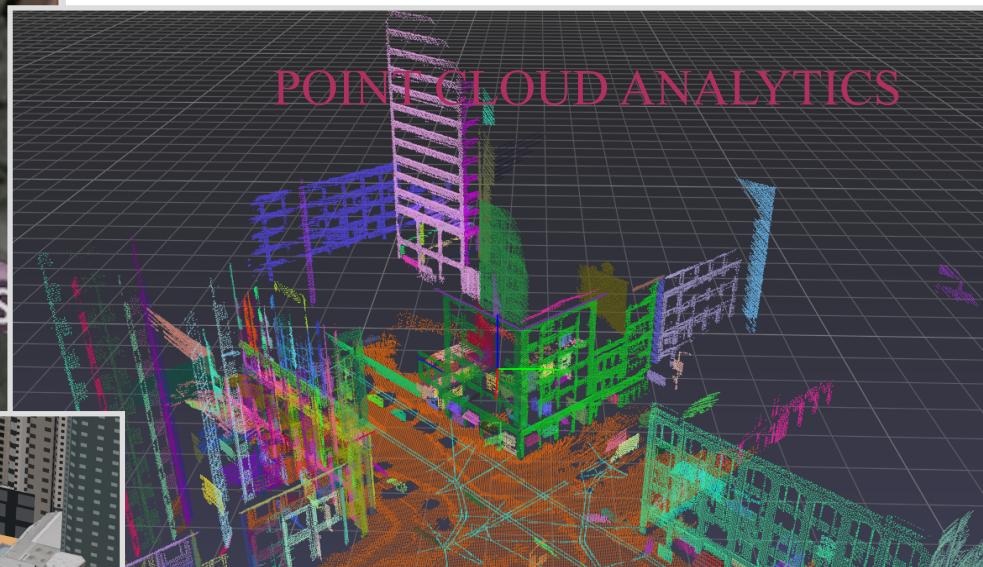
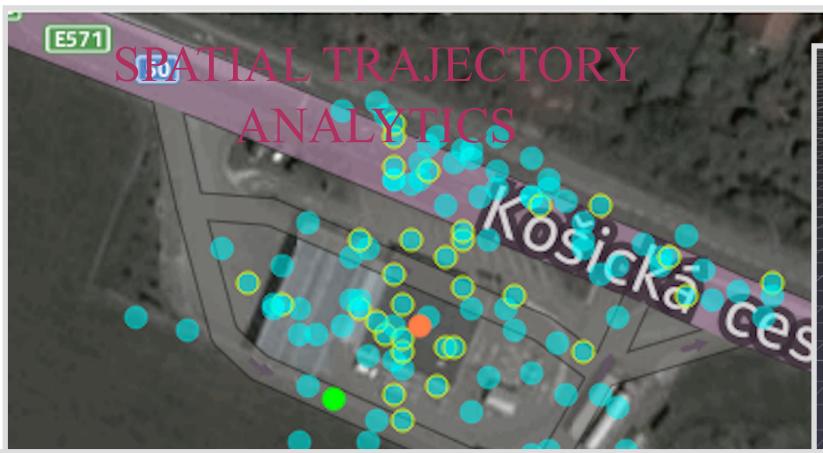
The power of location

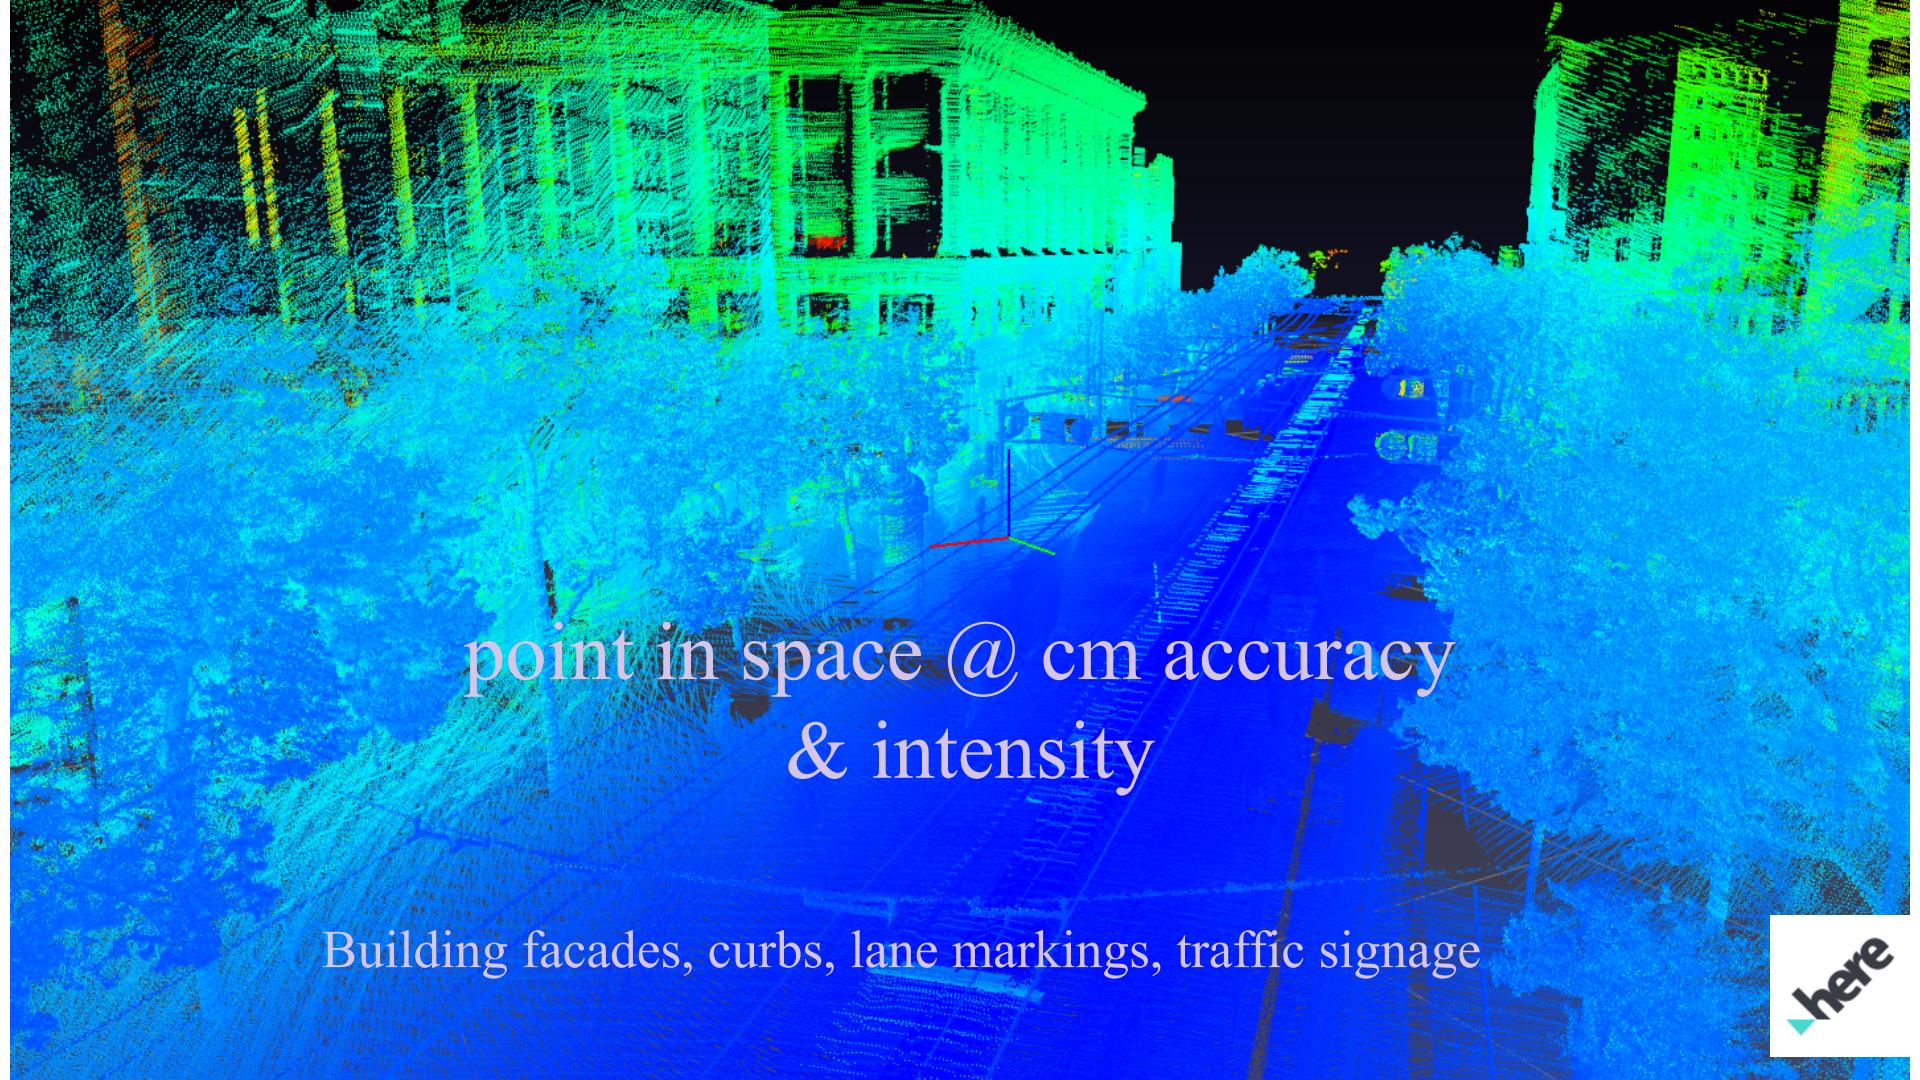


Research @ HERE



here



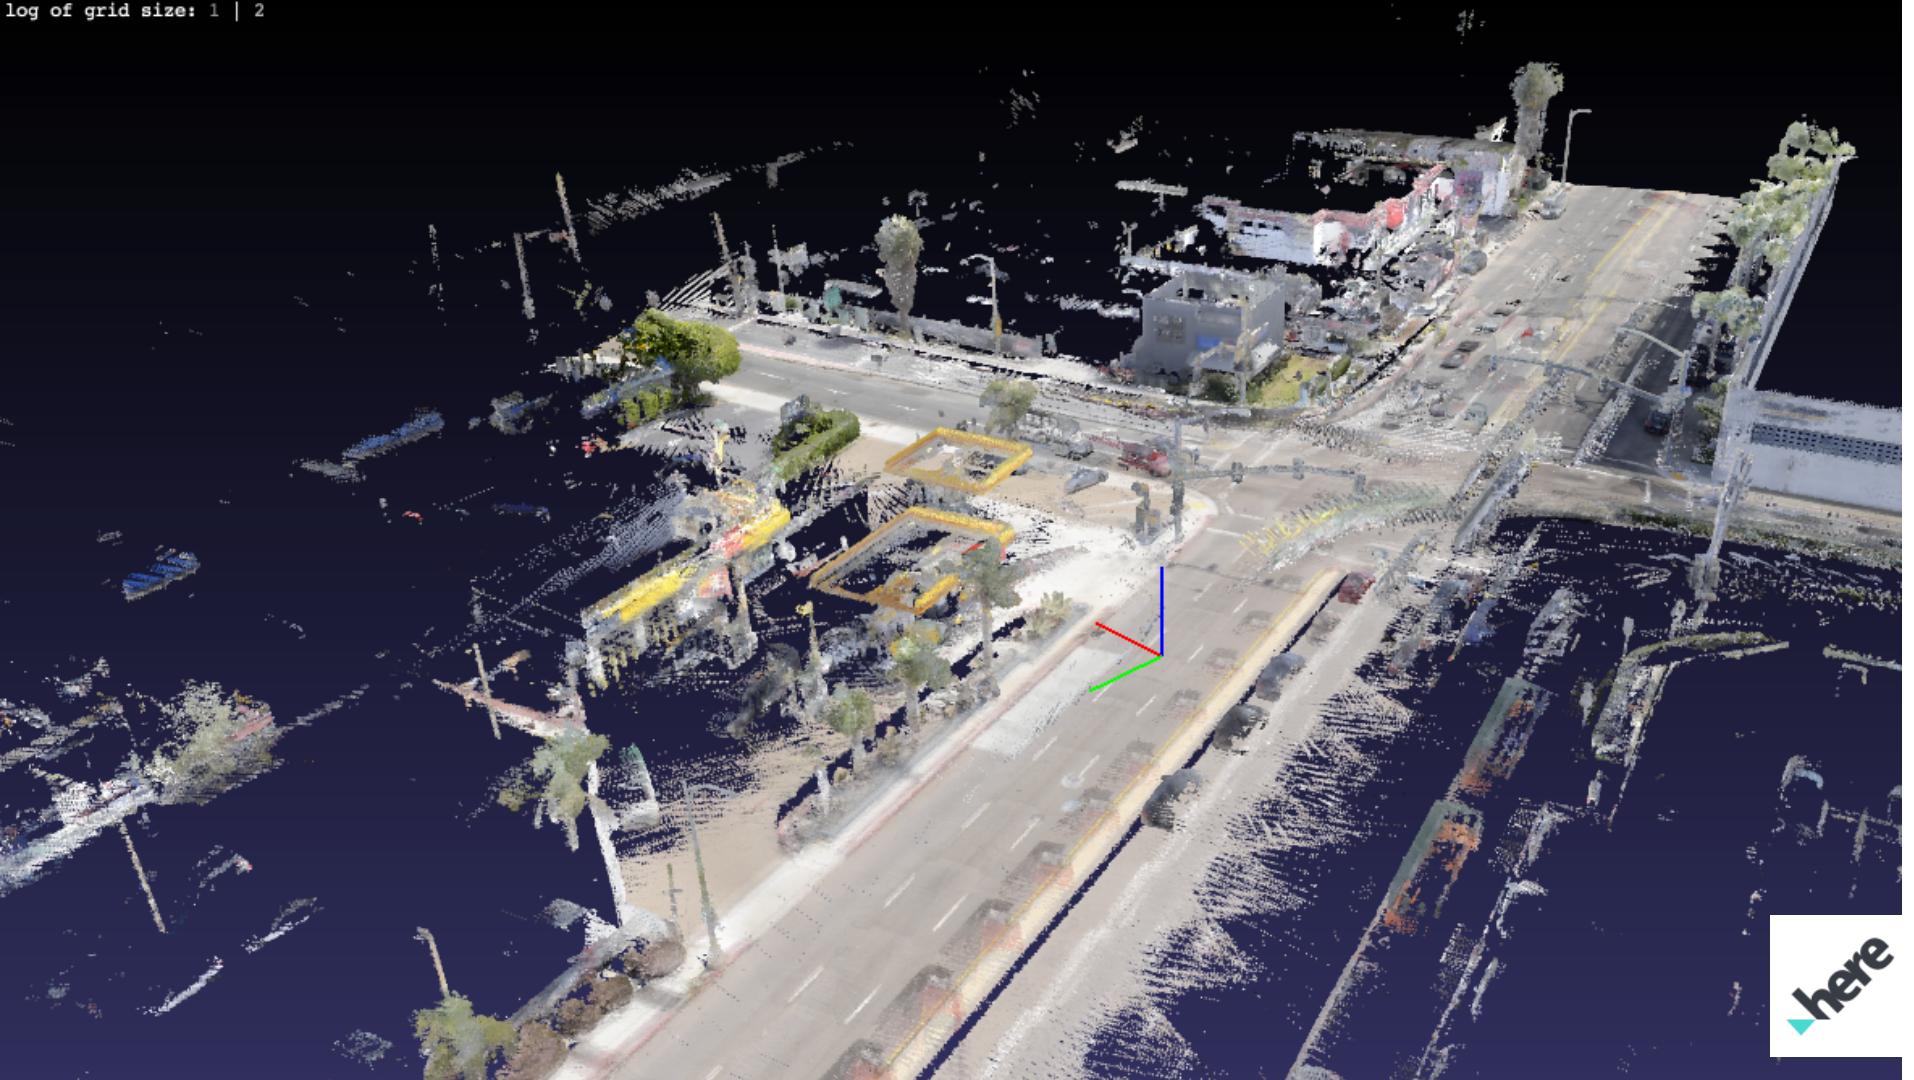


point in space @ cm accuracy
& intensity

Building facades, curbs, lane markings, traffic signage



log of grid size: 1 | 2



here

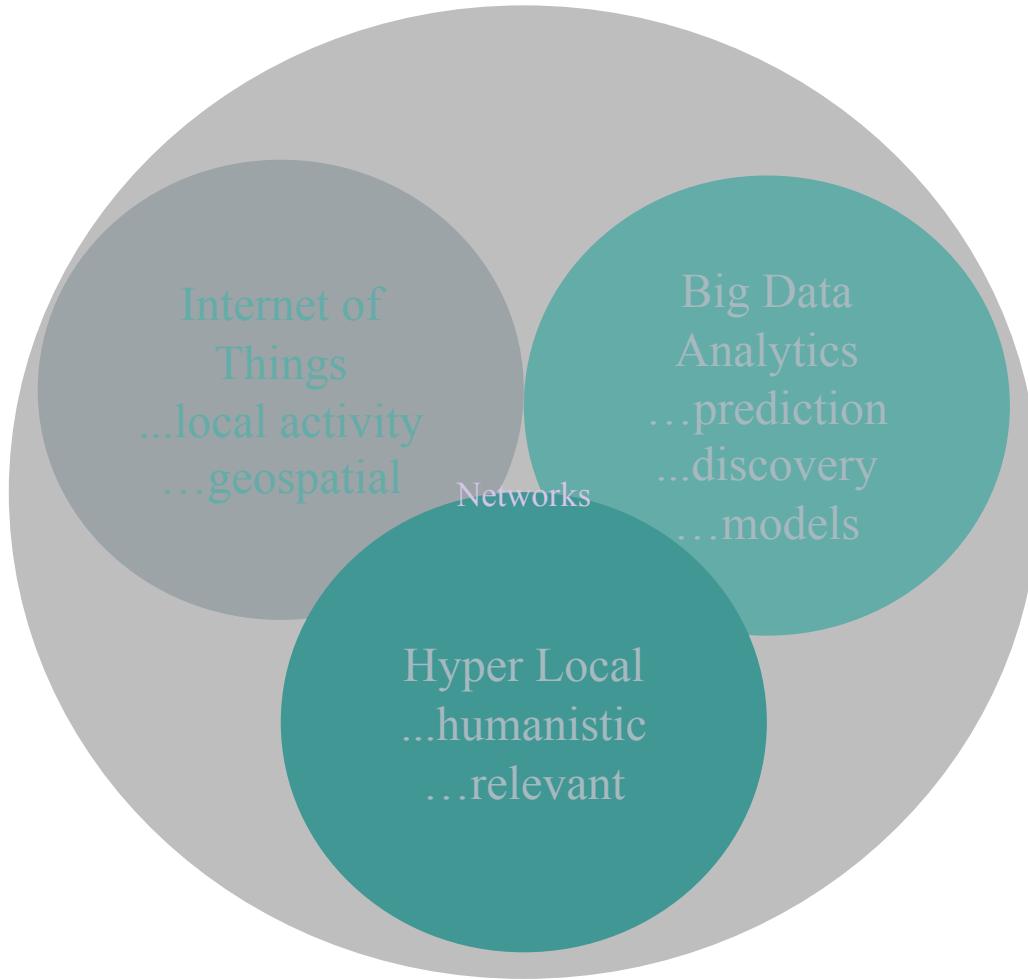
An aerial photograph of a park with a grid-like pattern of green grass and paved walkways. Numerous people are scattered across the park, some sitting on the grass in small groups, others walking along the paths. The scene is bright and suggests a sunny day.

Pervasive, data-driven,
Cloud-enabled
'Internet of Things'.

We are using multiple devices everyday

Machine generated
Human generated
Structured
Unstructured

Velocity
Volume
Variety
Variability
Veracity
Visualization
Value \$\$\$



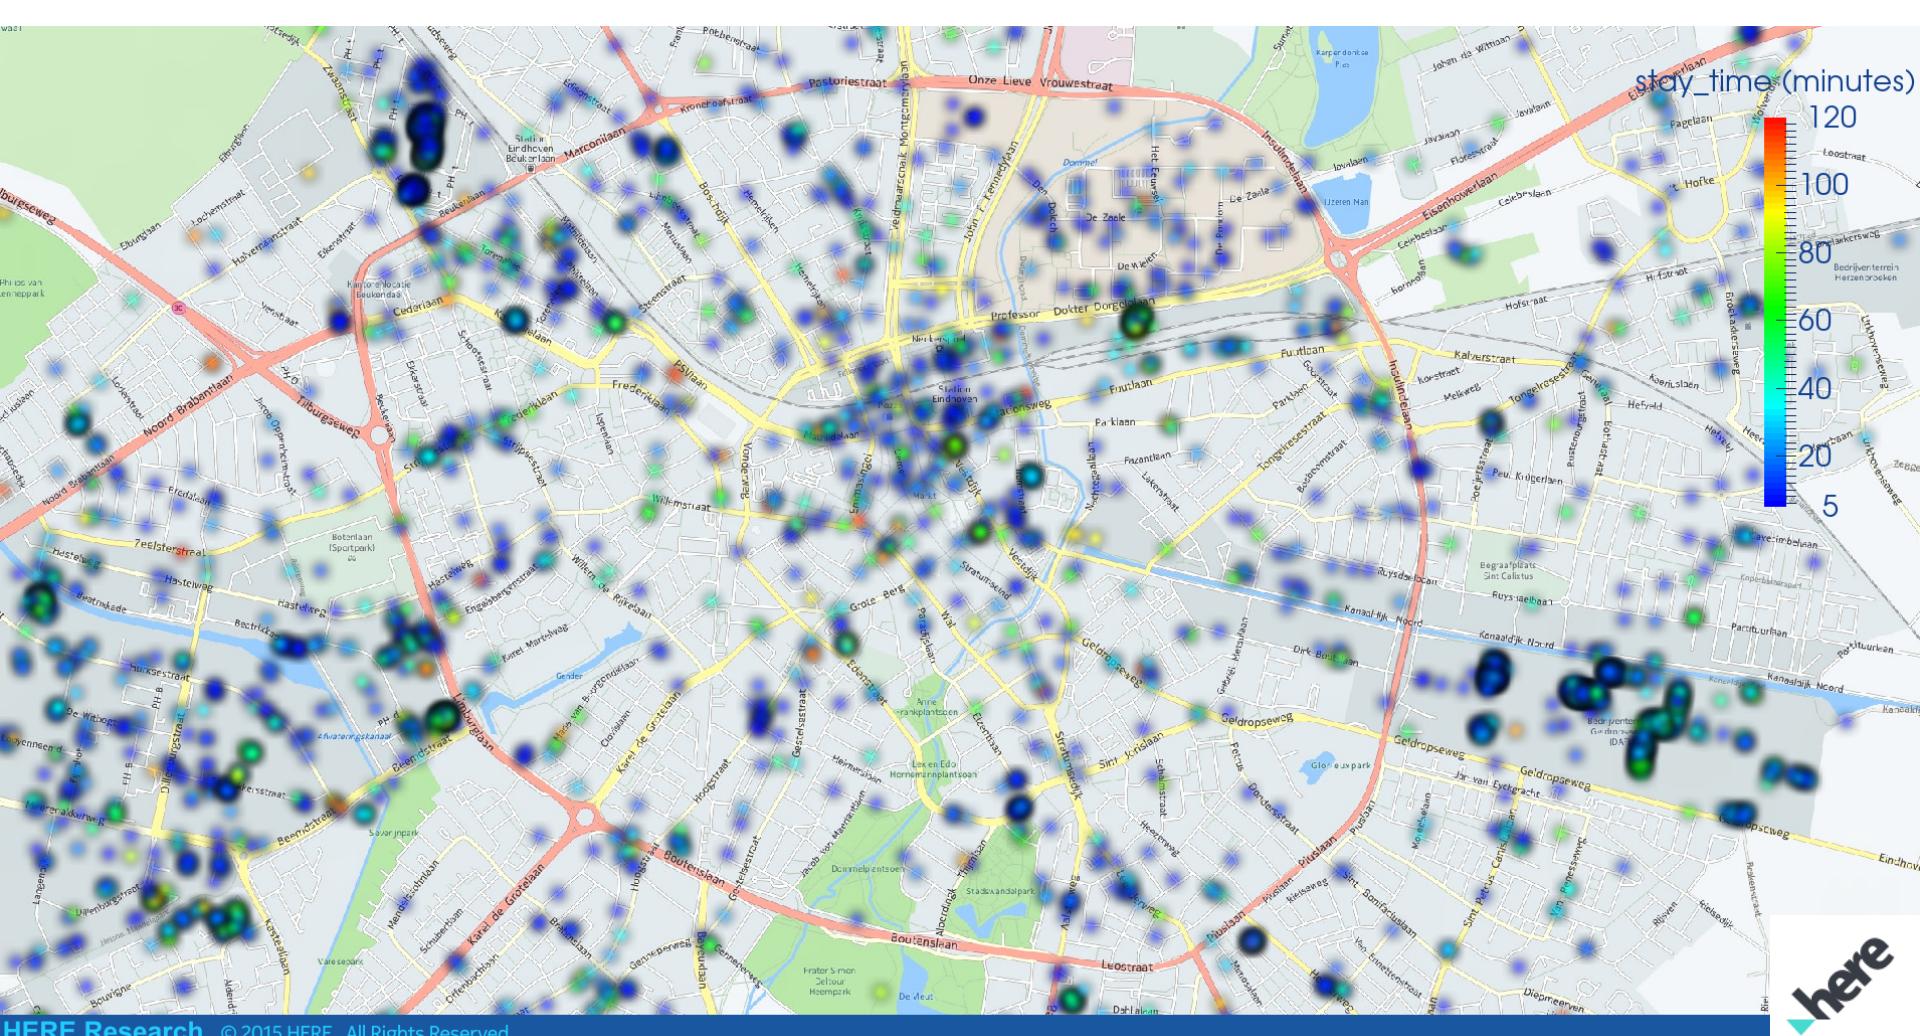
Maps are morphing into something new...

...companions, advisors
....assistants

....devices are listening
...reporting, tracking

here







Initial Speed
50-59kph

Centroid

Point of
first brake
application

Last second braking

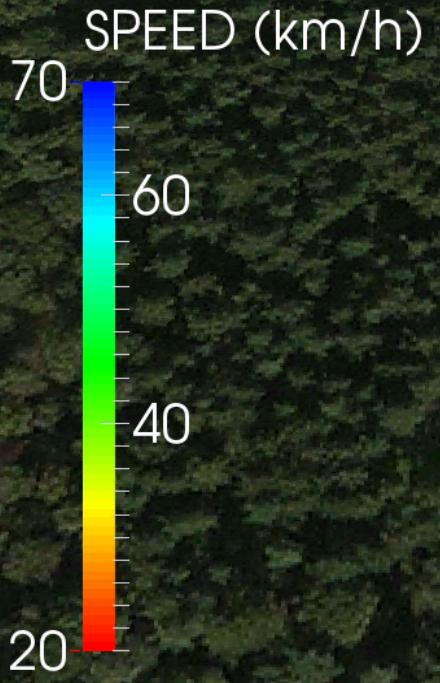
Normal Behavior
First Brake
Location

SPEED (km/h)



20

here



Abnormal Behavior
First Brake
Location



Lateral acceleration northbound direction

Speed [km/h]

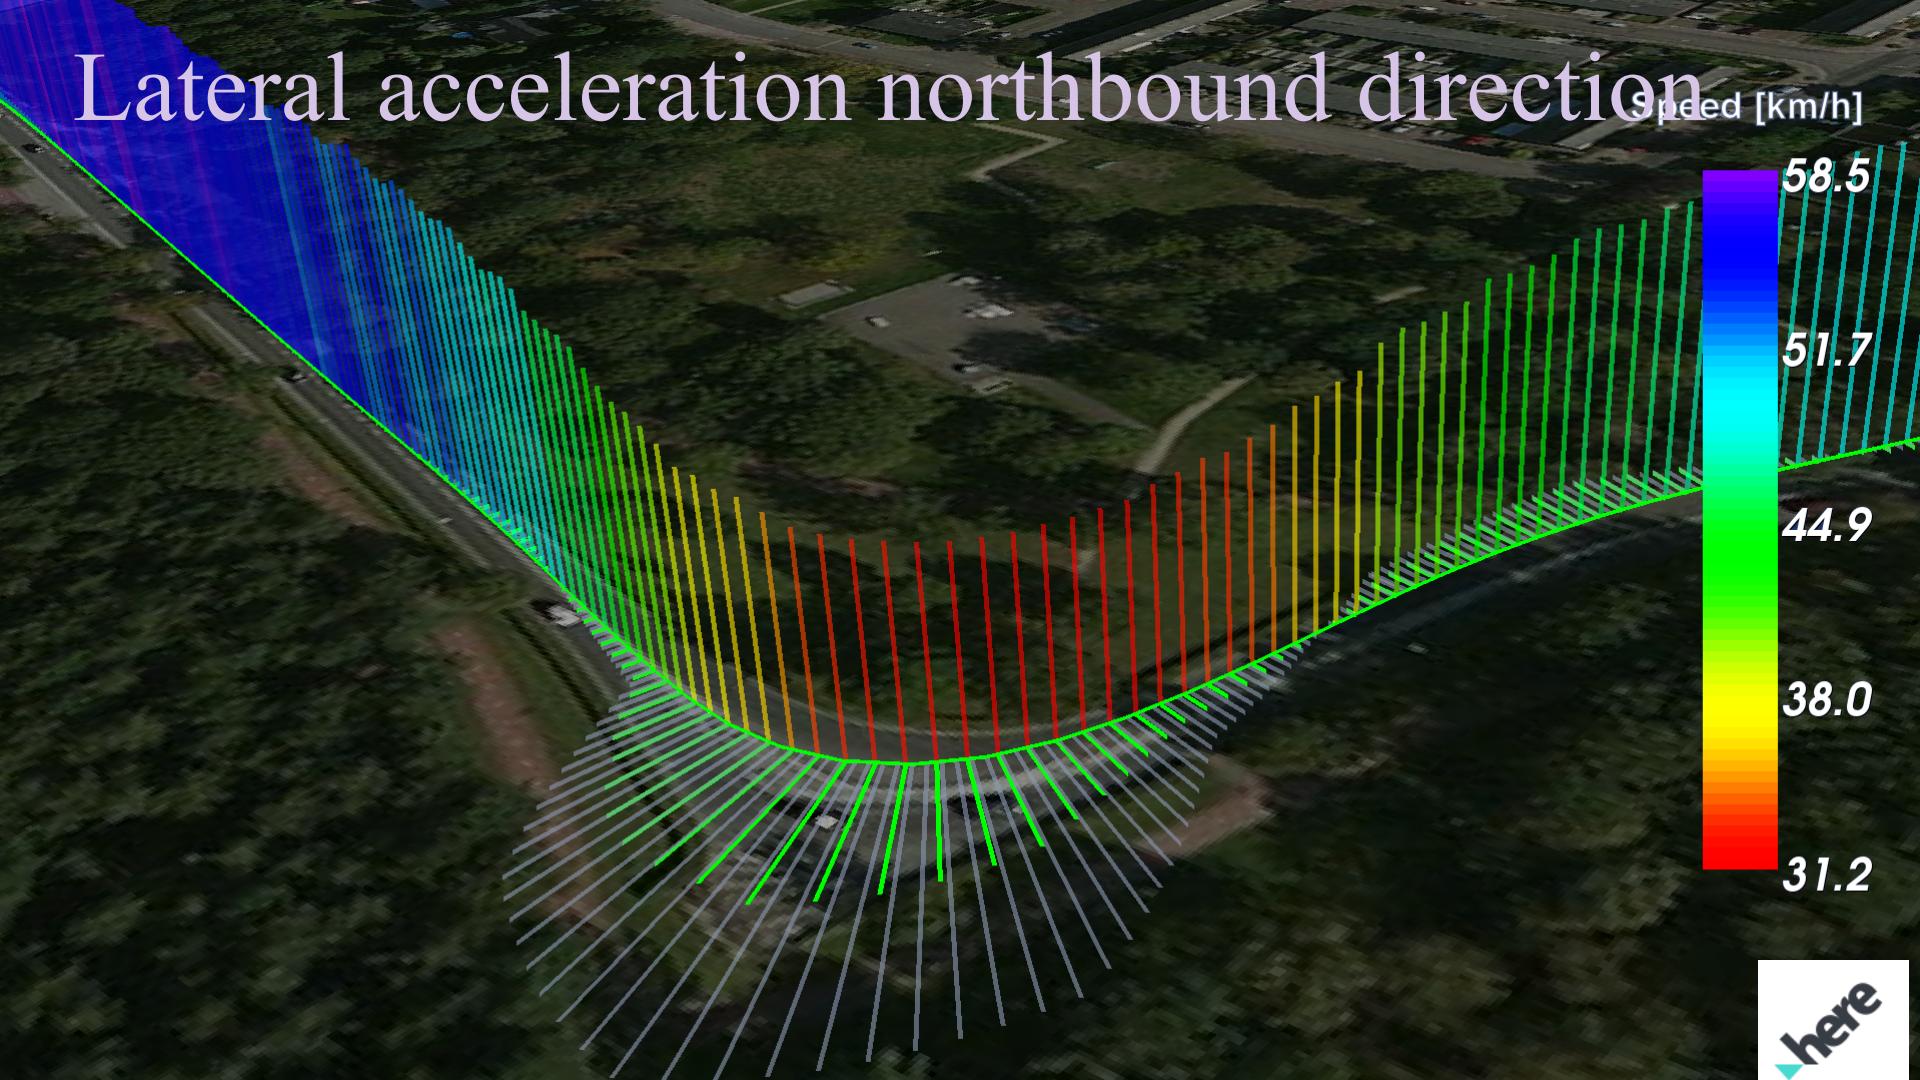
58.5

51.7

44.9

38.0

31.2



Lateral acceleration during rainstorm

Speed [km/h]

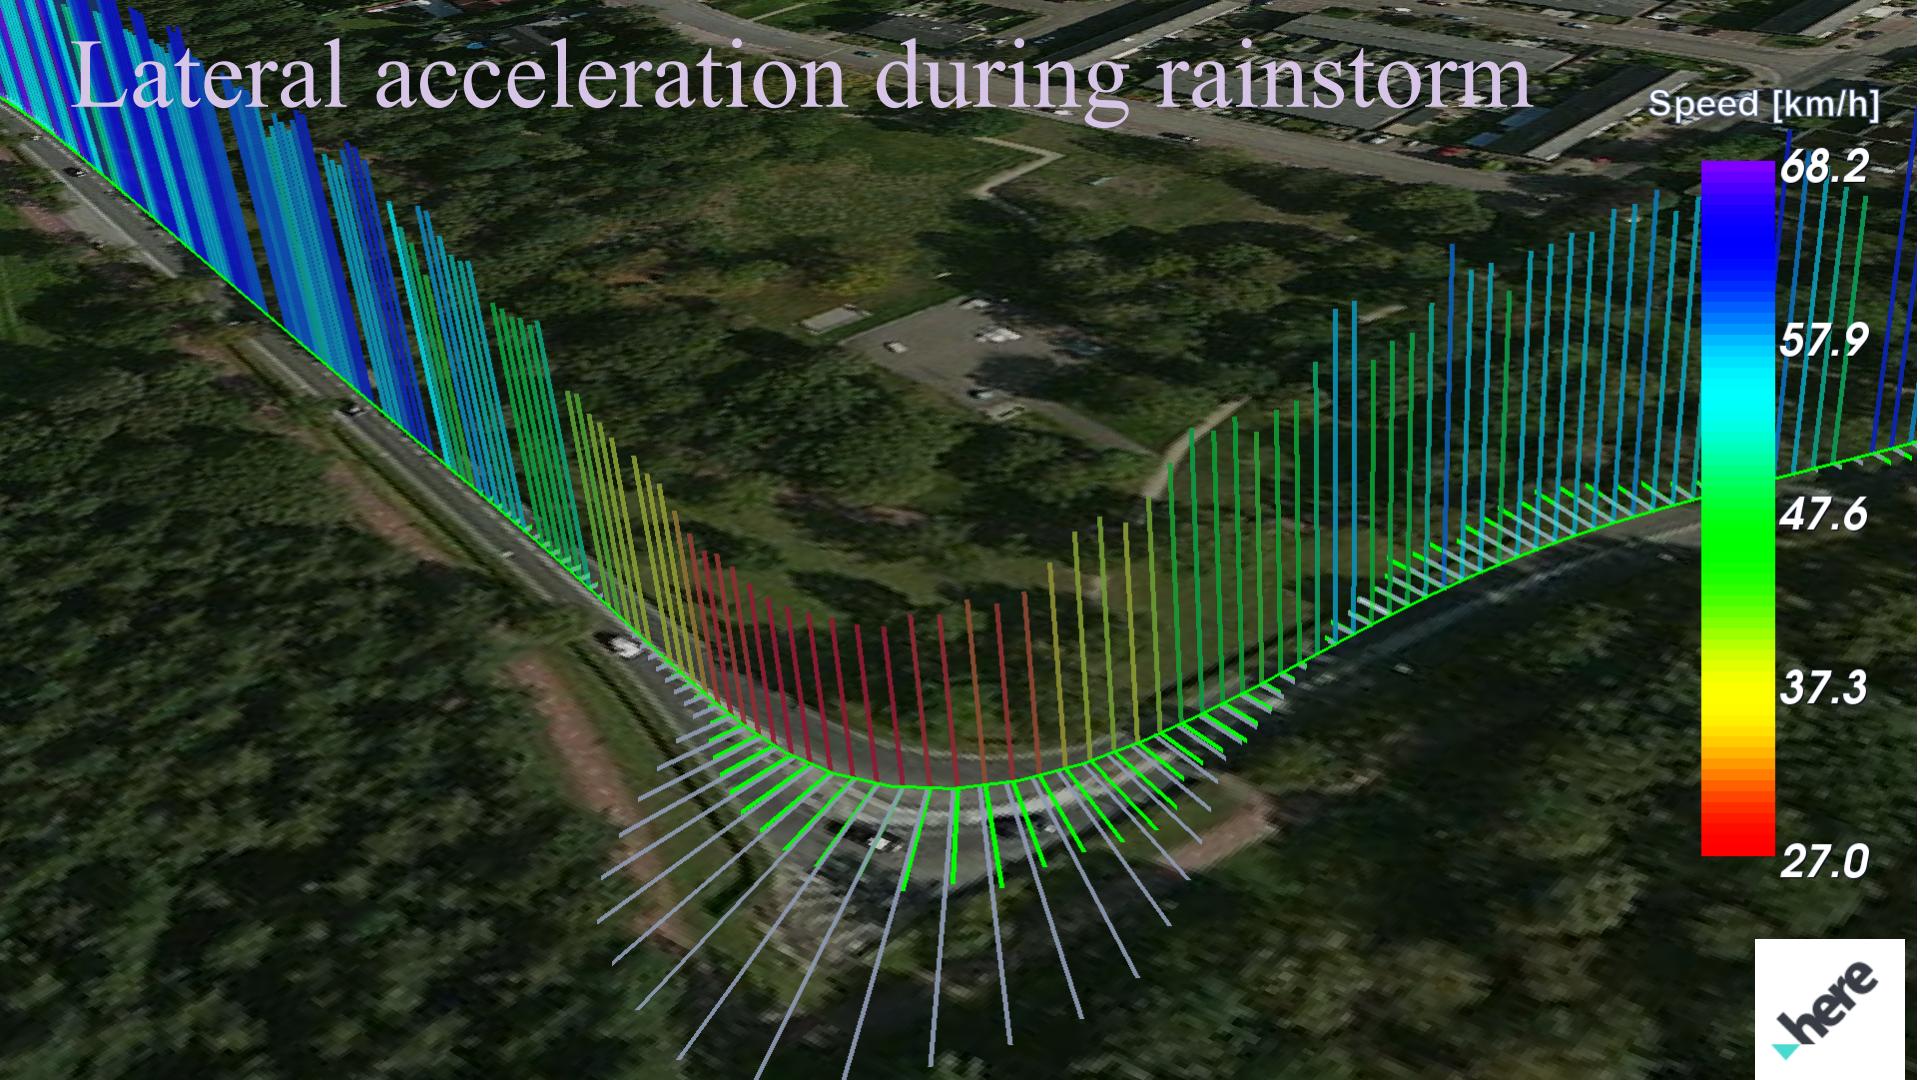
68.2

57.9

47.6

37.3

27.0



Use Cases with Deep Learning

here

Image classification

Objective:

Vehicle type identification based on images

Application:

Object identification for autonomous vehicles

Framework:

Image_Recognition_with_Keras

Sequence Prediction

Objective:

Predict number of rental bike over a time window

Application:

- Managing resources (bikes) to meet future demands
- Update points of interest of a map when someone is looking to rent a bike (e.g. likely to run out of available bikes in an hour)

Framework:

Sequence_Prediction_with_Keras



here